

## Your grade: 100%

Your latest: 100% • Your highest: 100% • To pass you need at least 80%. We keep your highest score.

[Next item →](#)

1. Suppose you learn a word embedding for a vocabulary of 10000 words. Then the embedding vectors could be 10000 dimensional, so as to capture the full range of variation and meaning in those words.

1 / 1 point

☐ True

☒ False

✓ **Correct**

The dimension of word vectors is usually smaller than the size of the vocabulary. Most common sizes for word vectors range between 50 and 1000.

2. What is t-SNE?

1 / 1 point

☐ A supervised learning algorithm for learning word embeddings

☐ A linear transformation that allows us to solve analogies on word vectors

☒ A non-linear dimensionality reduction technique

☐ An open-source sequence modeling library

✓ **Correct**

Yes

3. Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You then use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set.

x (input text)	y (happy?)
Having a great time!	1
I'm sad it's raining.	0
I'm feeling awesome!	1

Even if the word "wonderful" does not appear in your small training set, what label might be reasonably expected for the input text "I feel wonderful!"?

☐ y=0

☒ y=1

✓ **Correct**

Yes, word vectors empower your model with an incredible ability to generalize. The vector for "wonderful" would contain a negative/unhappy connotation which will probably make your model classify the sentence as a "1".

4. Which of these equations do you think should hold for a good word embedding? (Check all that apply)

☒  $e_{boy} - e_{girl} \approx e_{brother} - e_{sister}$

✔ **Correct**  
Yes!

☒  $e_{boy} - e_{brother} \approx e_{girl} - e_{sister}$

✔ **Correct**  
Yes!

☐  $e_{boy} - e_{girl} \approx e_{sister} - e_{brother}$

☐  $e_{boy} - e_{brother} \approx e_{sister} - e_{girl}$

5. Let  $A$  be an embedding matrix, and let  $o_{4567}$  be a one-hot vector corresponding to word 4567. Then to get the embedding of word 4567, why don't we call  $A * o_{4567}$  in Python?

- ☒ It is computationally wasteful.
- ☐ This doesn't handle unknown words (<UNK>).
- ☐ The correct formula is  $A^T * o_{4567}$ .
- ☐ None of the answers are correct: calling the Python snippet as described above is fine.

✔ **Correct**  
Yes, the element-wise multiplication will be extremely inefficient.

6. When learning word embeddings, we pick a given word and try to predict its surrounding words or vice versa.

- ☒ True
- ☐ False

✔ **Correct**  
Word embeddings are learned by picking a given word and trying to predict its surrounding words or vice versa.

7. True/False: In the word2vec algorithm, you estimate  $P(t | c)$ , where  $t$  is the target word and  $c$  is a context word.  $t$  and  $c$  are chosen from the training set to be nearby words.

- ☐ False
- ☒ True

✔ **Correct**  
Yes,  $t$  and  $c$  are chosen from the training set to be nearby words.

8. Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The word2vec model uses the following softmax function:

$$P(t | c) = \frac{e^{\theta_t^T e_c}}{\sum_{t'=1}^{10000} e^{\theta_{t'}^T e_c}}$$

Which of these statements are correct? Check all that apply.

- ☐  $\theta_t$  and  $e_c$  are both 10000 dimensional vectors.
- ☒  $\theta_t$  and  $e_c$  are both trained with an optimization algorithm such as Adam or gradient descent.

✔ Correct

- ☒  $\theta_t$  and  $e_c$  are both 500 dimensional vectors.

✔ Correct

- ☐ After training, we should expect  $\theta_t$  to be very close to  $e_c$  when  $t$  and  $c$  are the same word.

9. Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The GloVe model minimizes this objective:

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij})(\theta_i^T e_j + b_i + b_j' - \log X_{ij})^2$$

True/False:  $\theta_i$  and  $e_j$  should be initialized to 0 at the beginning of training.

- ☐ True
- ☒ False

✔ Correct

$\theta_i$  and  $e_j$  should be initialized randomly at the beginning of training.

10. You have trained word embeddings using a text dataset of  $s_1$  words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of  $s_2$  words. Keeping in mind that using word embeddings is a form of transfer learning, under which of these circumstances would you expect the word embeddings to be helpful?

- ☐  $s_1 \ll s_2$
- ☒  $s_1 \gg s_2$

✔ Correct

$s_1$  should transfer to  $s_2$